

Chapter 4


Network Layer: Data Plane

Ren Ping Liu

renping.liu@uts.edu.au

adapted from textbook slides by JFK/KWR

1 April 2024



*Computer Networking: A
Top-Down Approach*

8th Edition, Global Edition

Jim Kurose, Keith Ross

Copyright © 2022 Pearson Education Ltd

Network layer: “data plane” roadmap

4.1 Network layer: overview

- data plane
- control plane

4.2 What’s inside a router

- input ports, switching, output ports
- buffer management, scheduling

4.3 IP: the Internet Protocol

- datagram format
- addressing
- network address translation
- IPv6

~~4.4 Generalized Forwarding, SDN~~

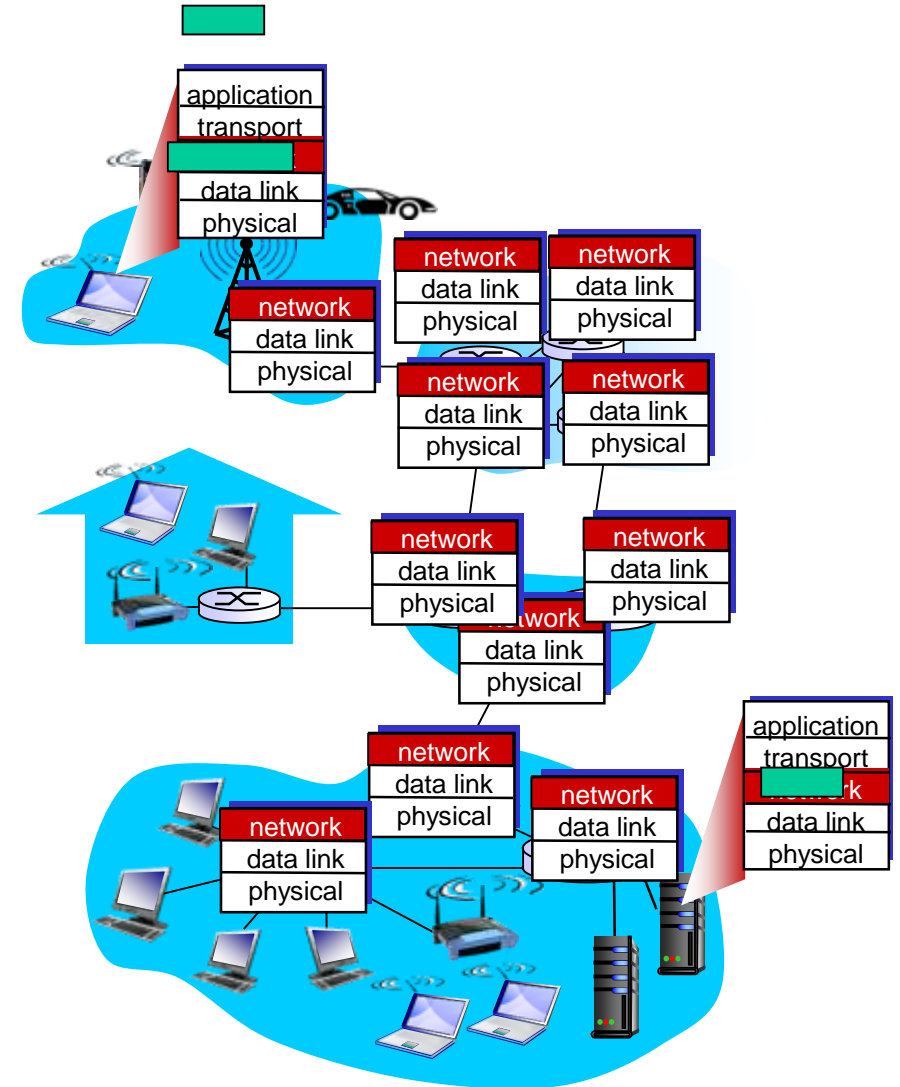
- ~~• Match+action~~
- ~~• OpenFlow: match+action in action~~

~~4.5 Middleboxes~~



Network-layer services and protocols

- transport segment from sending to receiving host
 - **sender:** encapsulates segments into datagrams, passes to link layer
 - **receiver:** delivers segments to transport layer protocol
- network layer protocols in *every Internet device*: hosts, routers
- **routers:**
 - examines header fields in all IP datagrams passing through it
 - moves datagrams from input ports to output ports to transfer datagrams along end-end path



Two key network-layer functions

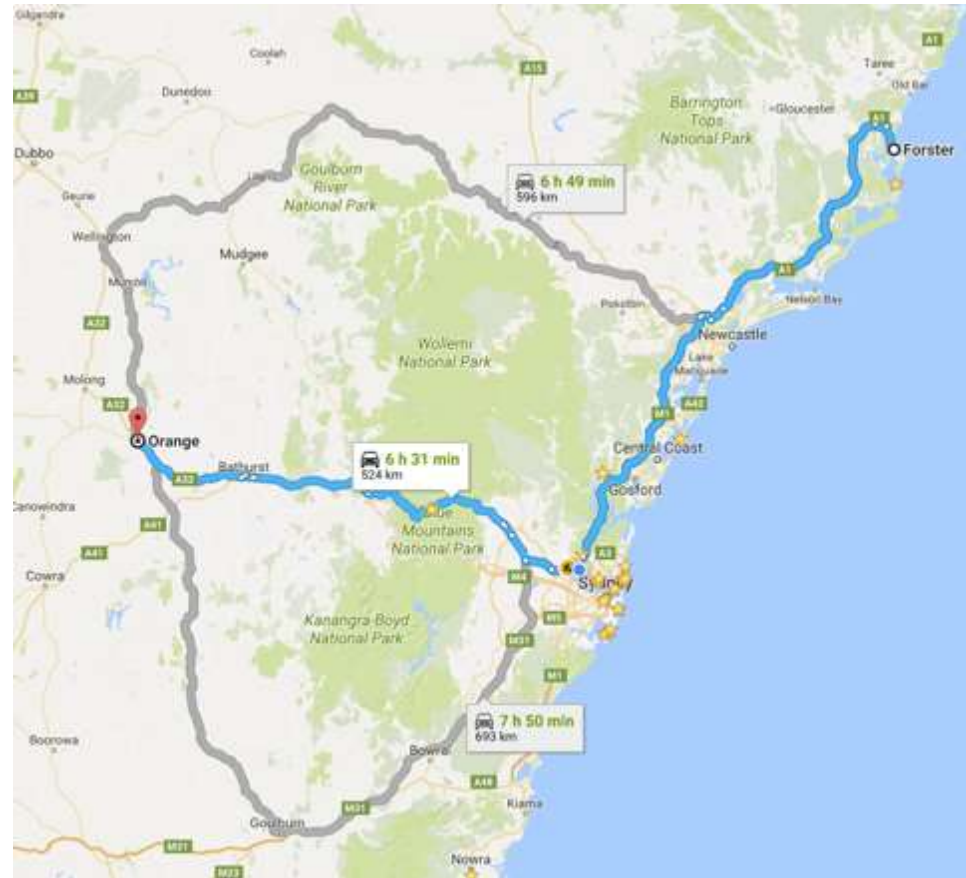
Trip: Forster → Orange

network-layer functions:

- *routing*: determine route taken by packets from source to destination

- *routing algorithms*


- *forwarding*: move packets from router's input to appropriate router output



Network layer: data plane, control plane


Control plane

- determines how datagram is routed among routers along end-end path from source host to destination host
- network-wide logic
- Based on network topology



Data plane

- determines how datagram arriving on router input port is forwarded to router output port
- local, per-router function
- Based on pre-determined forwarding



Network layer: data plane, control plane

Control plane

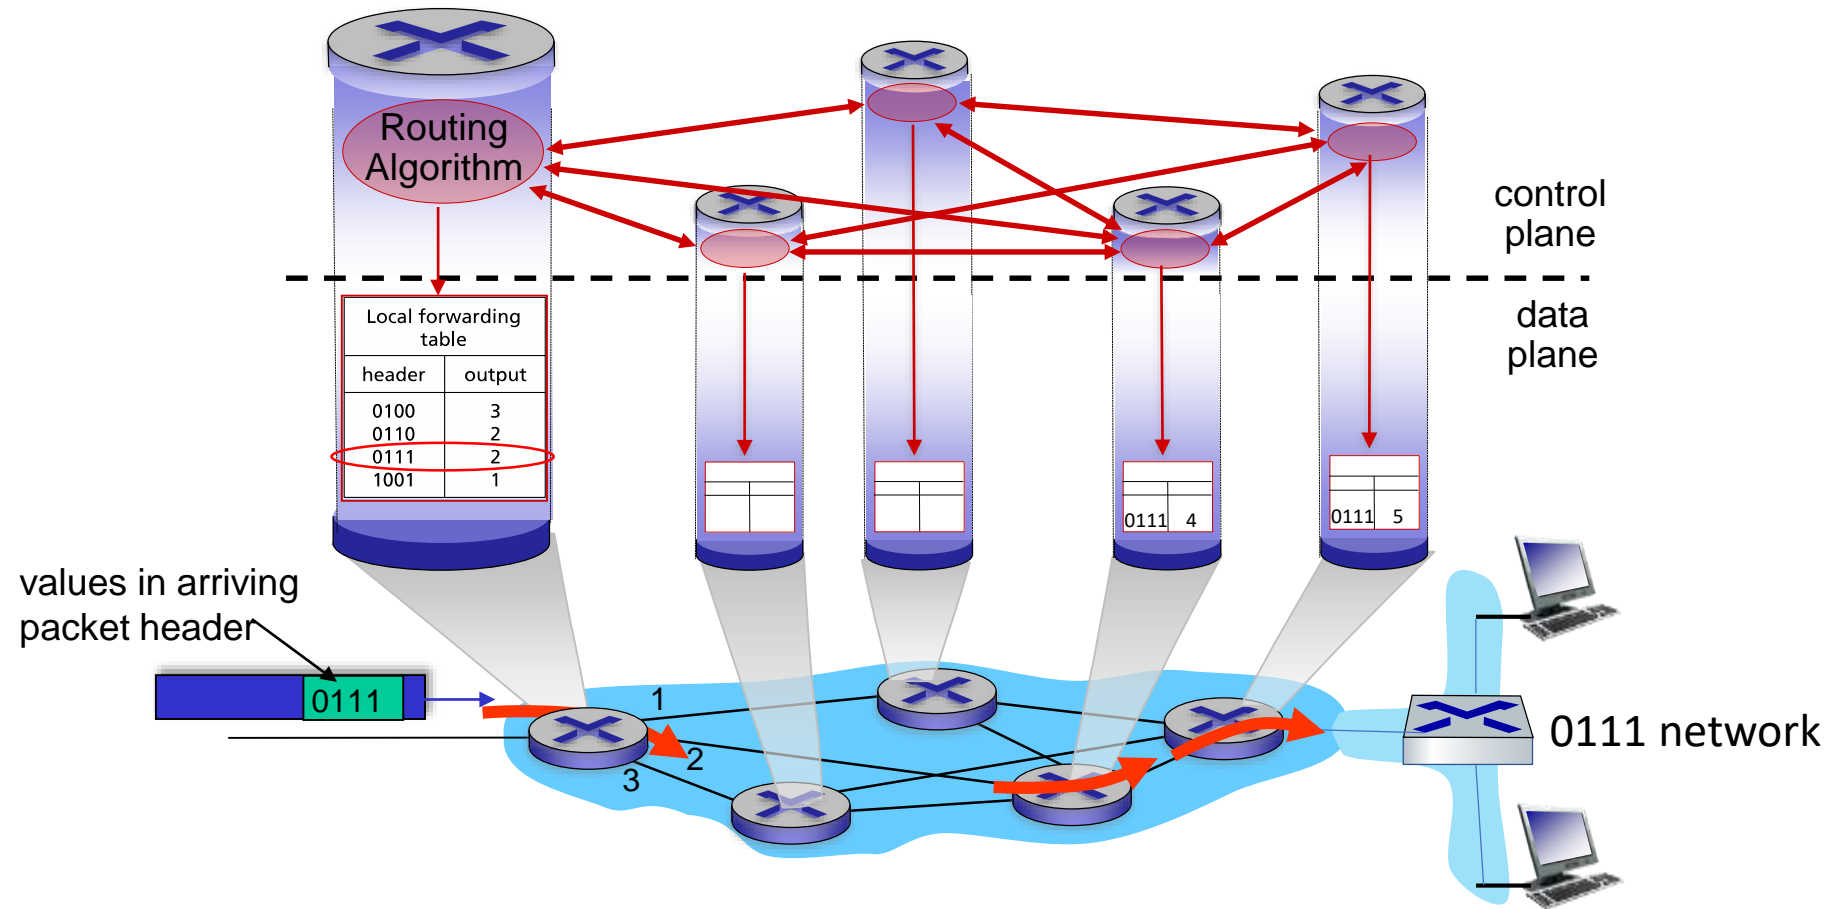
- determines how datagram is routed among routers along end-end path from source host to destination host
- network-wide logic
- Based on network topology

Two control-plane approaches:

- *traditional routing algorithms*: implemented in routers
- *software-defined networking (SDN)*: implemented in (remote) servers

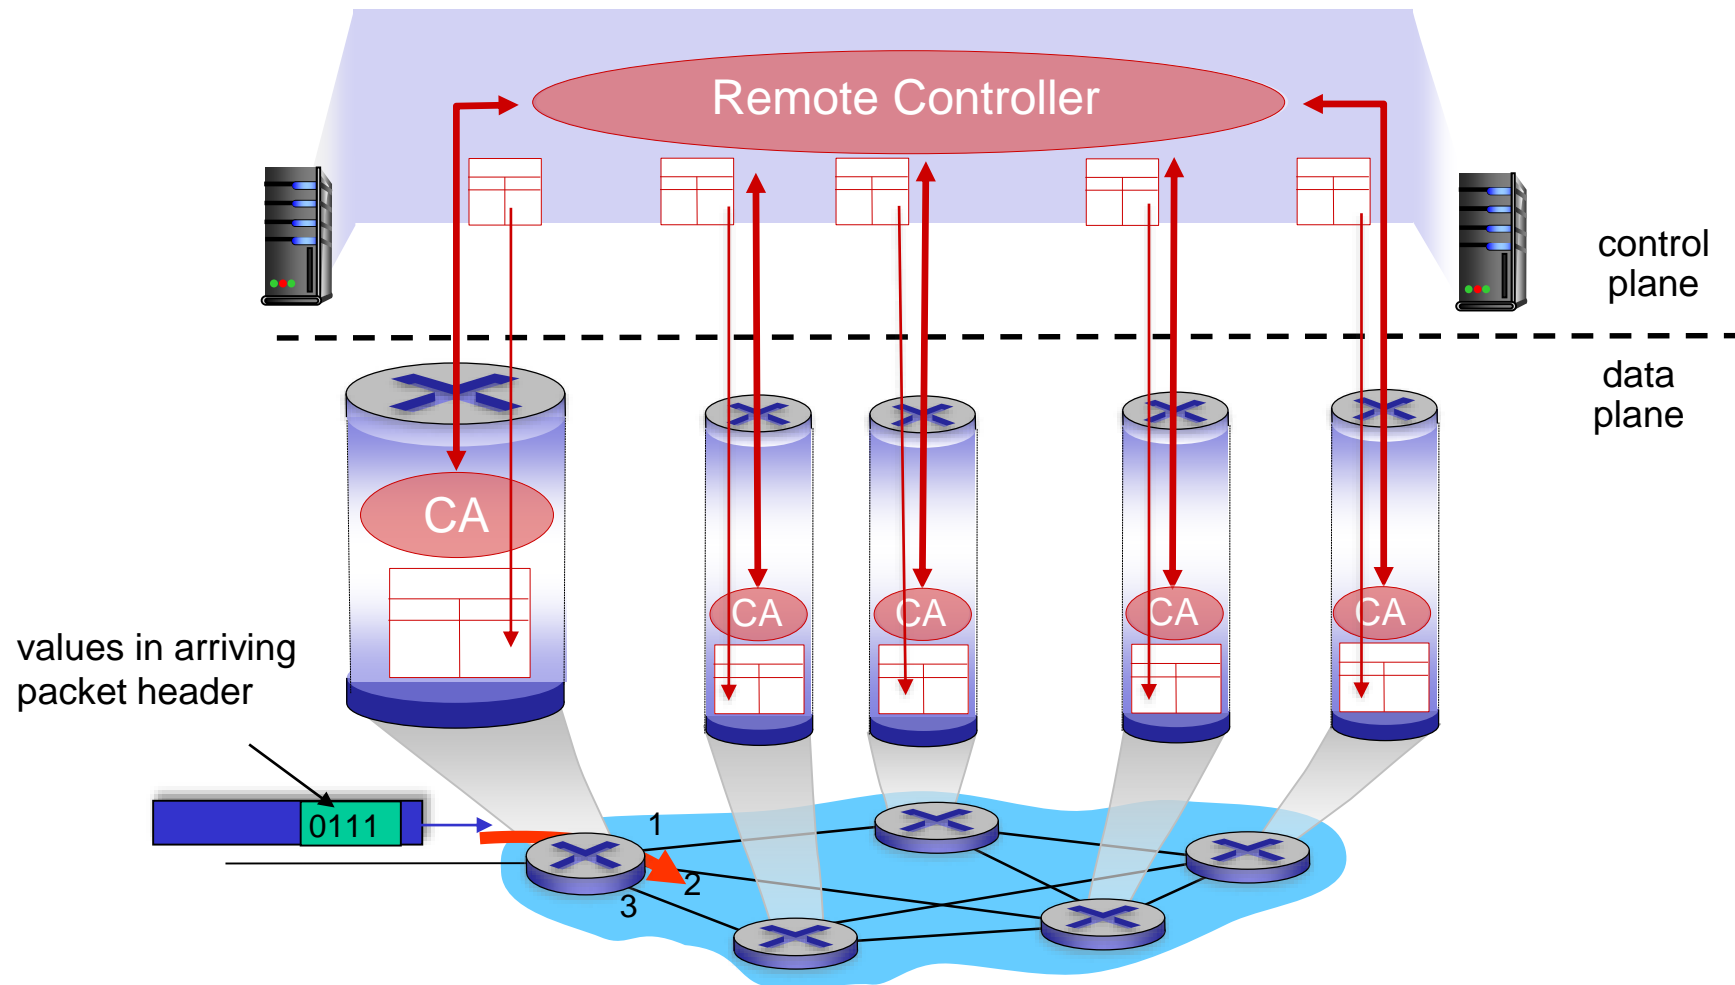
Per-router control plane

Individual routing algorithm components *in each and every router* interact in the control plane



Software-Defined Networking (SDN) control plane

Remote controller computes, installs forwarding tables in routers



Network service model

think of services by Aust Post: letter, registered, express, parcel,

Q: What *service model* do we need for “network” to pass datagrams from sender to receiver?

example services for
individual datagrams:

- guaranteed delivery
- guaranteed timing, e.g., delivery within 40 msec delay

example services for a *flow* of datagrams:

- in-order datagram delivery
- guaranteed minimum bandwidth to flow
- restrictions on changes in inter-packet spacing

Network-layer service model

Network Architecture	Service Model	Quality of Service (QoS) Guarantees ?			
		Bandwidth	Loss	Order	Timing
Internet	best effort	none	no	no	no

Internet “best effort” service model

No guarantees on:

- ✗ successful datagram delivery to destination
- ✗ timing or order of delivery
- ✗ bandwidth available to end-end flow

Network-layer service model

Network Architecture	Service Model	Quality of Service (QoS) Guarantees ?			
		Bandwidth	Loss	Order	Timing
Internet	best effort	none	no	no	no

QoS: too hard, too expensive!

Reflections on best-effort service:

- **simplicity of mechanism** has allowed Internet to be widely deployed adopted
- sufficient **provisioning of bandwidth** allows performance of real-time applications (e.g., interactive voice, video) to be “good enough” for “most of the time”
- **replicated, application-layer distributed services** (datacenters, content distribution networks) connecting close to clients’ networks, allow services to be provided from multiple locations
- congestion control of “elastic” services helps

It's hard to argue with success of best-effort service model

Network layer: “data plane” roadmap

4.1 Network layer: overview

- data plane
- control plane

4.2 What’s inside a router

- input ports, switching, output ports
- buffer management, scheduling

4.3 IP: the Internet Protocol

- datagram format
- addressing
- network address translation
- IPv6



~~4.4 Generalized Forwarding, SDN~~

- ~~• Match+action~~
- ~~• OpenFlow: match+action in action~~

~~4.5 Middleboxes~~

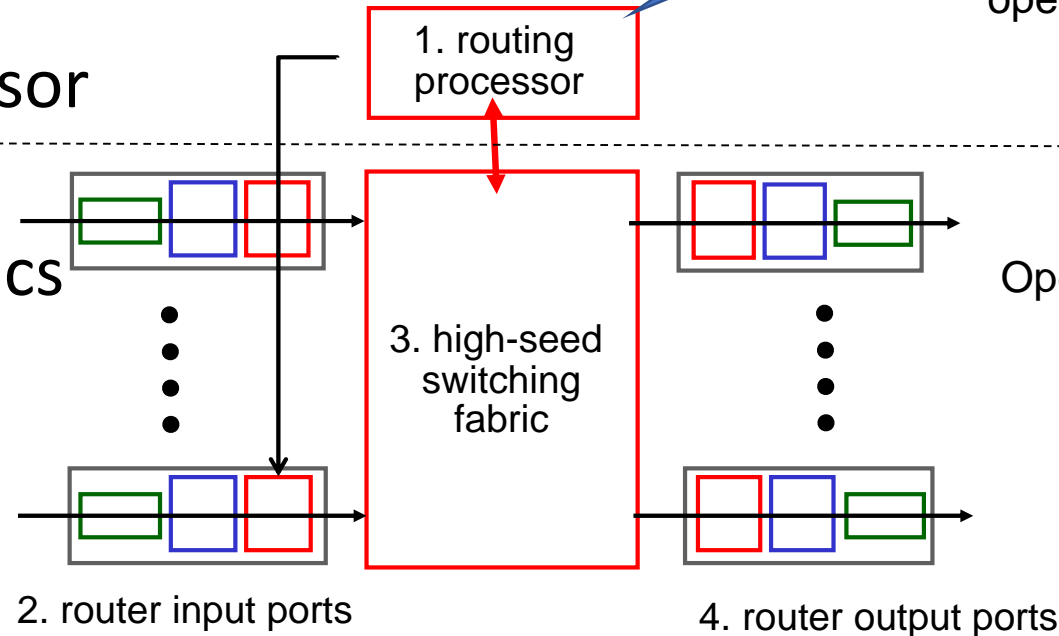
Router architecture overview



A specialised computer

■ Four parts:


- | | |
|----------|----------------------|
| Software | 1. routing processor |
| <hr/> | |
| Hardware | 2. Input ports |
| | 3. Switching fabrics |
| | 4. Output ports |



control plane (software)
operates in **ms** time frame
routing, management

data plane (hardware)
Operates in **ns** timeframe
forwarding datagram

Input port functions




physical layer:
bit-level reception


link layer:
e.g., Ethernet
(chapter 6)

decentralized switching:

- using header field values, lookup output port using forwarding table in input port memory (*“match plus action”*)
- goal: complete input port processing at ‘line speed’
- **input port queuing:** if datagrams arrive faster than forwarding rate into switch fabric



Input port functions




physical layer:
bit-level reception

link layer:
e.g., Ethernet
(chapter 6)


decentralized switching:

- using header field values, lookup output port using forwarding table in input port memory (*"match plus action"*)
- **destination-based forwarding:** forward based only on destination IP address (traditional)
- ~~**generalized forwarding:** forward based on any set of header field values~~




Input port queuing

- If switch fabric slower than input ports combined -> queueing may occur at input queues
 - queueing delay and loss due to input buffer overflow!
- **Head-of-the-Line (HOL) blocking:** queued datagram at front of queue prevents others in queue from moving forward




output port contention: only one red datagram can be transferred. lower red packet is *blocked*



one packet time later: green packet experiences HOL blocking




Switching fabrics

- transfer packet from input link to appropriate output link
- **switching rate**: rate at which packets can be transfer from inputs to outputs
 - often measured as multiple of input/output line rate
 - N inputs: switching rate N times line rate desirable



Switching fabrics


- transfer packet from input link to appropriate output link
- **switching rate**: rate at which packets can be transfer from inputs to outputs
 - often measured as multiple of input/output line rate
 - N inputs: switching rate N times line rate desirable
- three major types of switching fabrics:



Switching via memory


first generation routers:

- traditional computers with switching under direct control of CPU
- packet copied to system's memory
- speed limited by memory bandwidth (2 bus crossings per datagram)





Switching via a bus

- datagram from input port memory to output port memory via a shared bus
- *bus contention*: switching speed limited by bus bandwidth
- 32 Gbps bus, Cisco 5600: sufficient speed for access routers

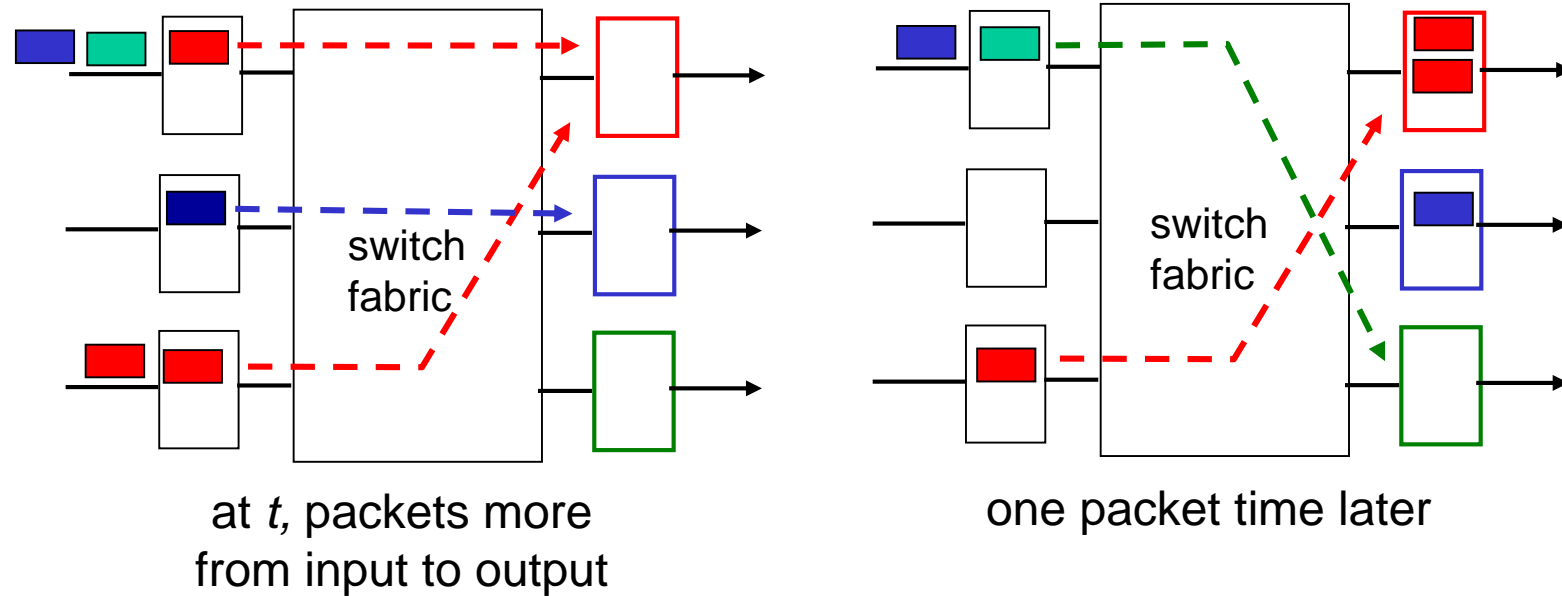


Switching via interconnection network

- Crossbar, Clos networks, other interconnection nets initially developed to connect processors in multiprocessor
- **multistage switch**: $n \times n$ switch from multiple stages of smaller switches
- **exploiting parallelism**:
 - fragment datagram into fixed length cells on entry
 - switch cells through the fabric, reassemble datagram at exit

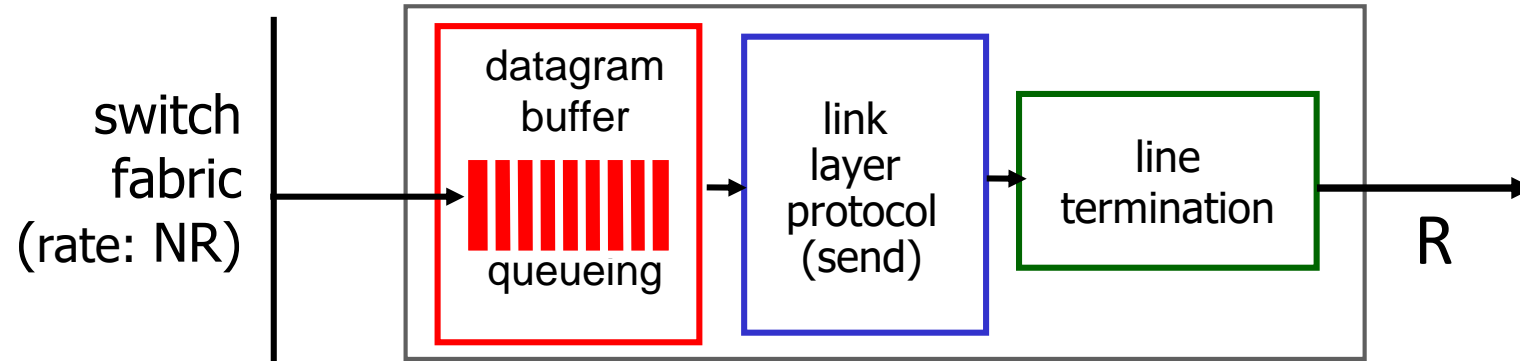


Output port queuing



- buffering when arrival rate via switch exceeds output line speed
- *queueing (delay) and loss due to output port buffer overflow!*

Output port queuing



- **Buffering** required when datagrams arrive from fabric faster than link transmission rate.
- **Scheduling discipline** chooses among queued datagrams for transmission → Priority scheduling – who gets best performance, network neutrality
- **Drop policy**: which datagrams to drop if no free buffers? → Datagrams can be lost due to congestion, lack of buffers

Packet Scheduling: FCFS


packet scheduling: deciding which packet to send next on link

- first come, first served
- priority
- round robin
- weighted fair queueing

FCFS: packets transmitted in order of arrival to output port

- also known as: First-in-first-out (FIFO)
- real world examples?

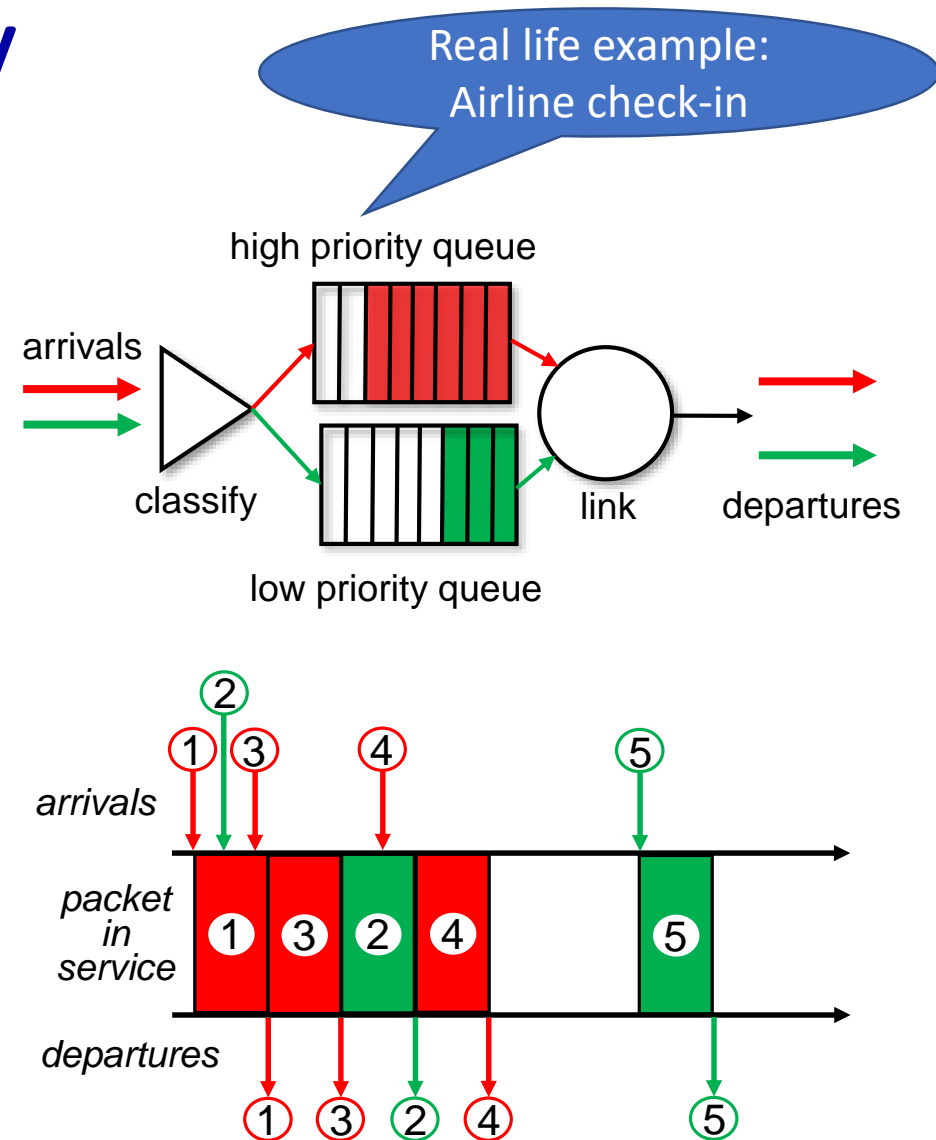
Abstraction: queue



Scheduling policies: priority

Priority scheduling:


- arriving traffic classified, queued by class
 - any header fields can be used for classification
- send packet from highest priority queue that has buffered packets
 - FCFS within priority class



Scheduling policies: round robin

Round Robin (RR) scheduling:

- arriving traffic classified, queued by class
 - any header fields can be used for classification
- server cyclically, repeatedly scans class queues, sending one complete packet from each class (if available) in turn




Scheduling policies: weighted fair queueing

Weighted Fair Queuing (WFQ):

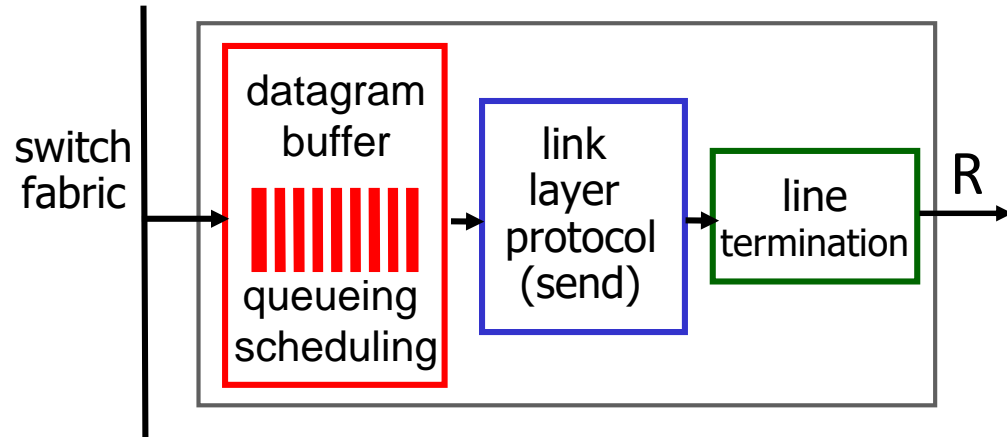
- generalized Round Robin
- each class, i , has weight, w_i , and gets weighted amount of service in each cycle:

$$\frac{w_i}{\sum_j w_j}$$


- minimum bandwidth guarantee (per-traffic-class)



Buffer Management



Abstraction: queue



buffer management:

- **drop:** which packet to add, drop when buffers are full
 - **tail drop:** drop arriving packet
 - **priority:** drop/remove on priority basis
- **marking:** which packets to mark to signal congestion (ECN, RED)

Sidebar: Network Neutrality

What is network neutrality?

- *technical*: how an ISP should share/allocation its resources
 - packet scheduling, buffer management are the *mechanisms*
- *social, economic* principles
 - protecting free speech
 - encouraging innovation, competition
- enforced *legal* rules and policies

Optional – not tested

Different countries have different “takes” on network neutrality

Sidebar: Network Neutrality

2015 US FCC *Order on Protecting and Promoting an Open Internet*: three “clear, bright line” rules:

- **no blocking** ... “shall not block lawful content, applications, services, or non-harmful devices, subject to reasonable network management.”
- **no throttling** ... “shall not impair or degrade lawful Internet traffic on the basis of Internet content, application, or service, or use of a non-harmful device, subject to reasonable network management.”
- **no paid prioritization.** ... “shall not engage in paid prioritization”

Optional

ISP: telecommunications or information service?

Is an ISP a “telecommunications service” or an “information service” provider?

- the answer *really* matters from a regulatory standpoint!

US Telecommunication Act of 1934 and 1996:

- *Title II*: imposes “common carrier duties” on *telecommunications services*: reasonable rates, non-discrimination and *requires regulation*
- *Title I*: applies to *information services*:
 - no common carrier duties (*not regulated*)
 - but grants FCC authority “... as may be necessary in the execution of its functions”

Optional

– not tested

Mid-break



■ Q & A



Network layer: “data plane” roadmap

4.1 Network layer: overview

- data plane
- control plane

4.2 What’s inside a router

- input ports, switching, output ports
- buffer management, scheduling

4.3 IP: the Internet Protocol

- datagram format
- addressing
- network address translation
- IPv6



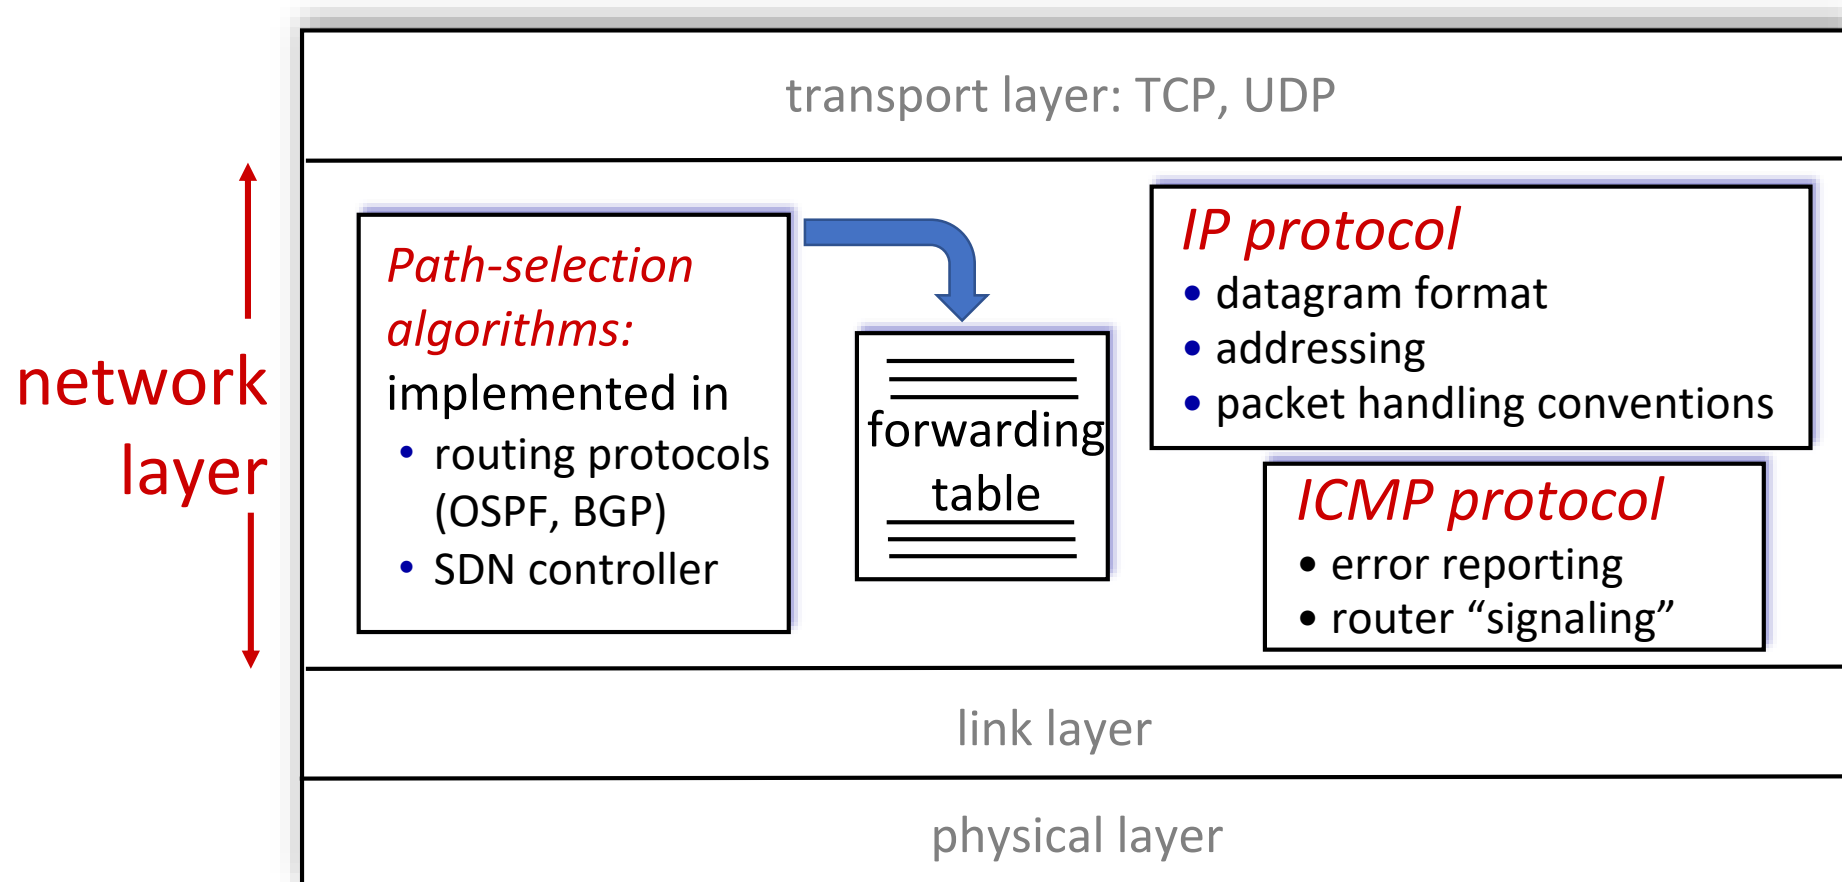
~~4.4 Generalized Forwarding, SDN~~

- ~~• Match+action~~
- ~~• OpenFlow: match+action in action~~

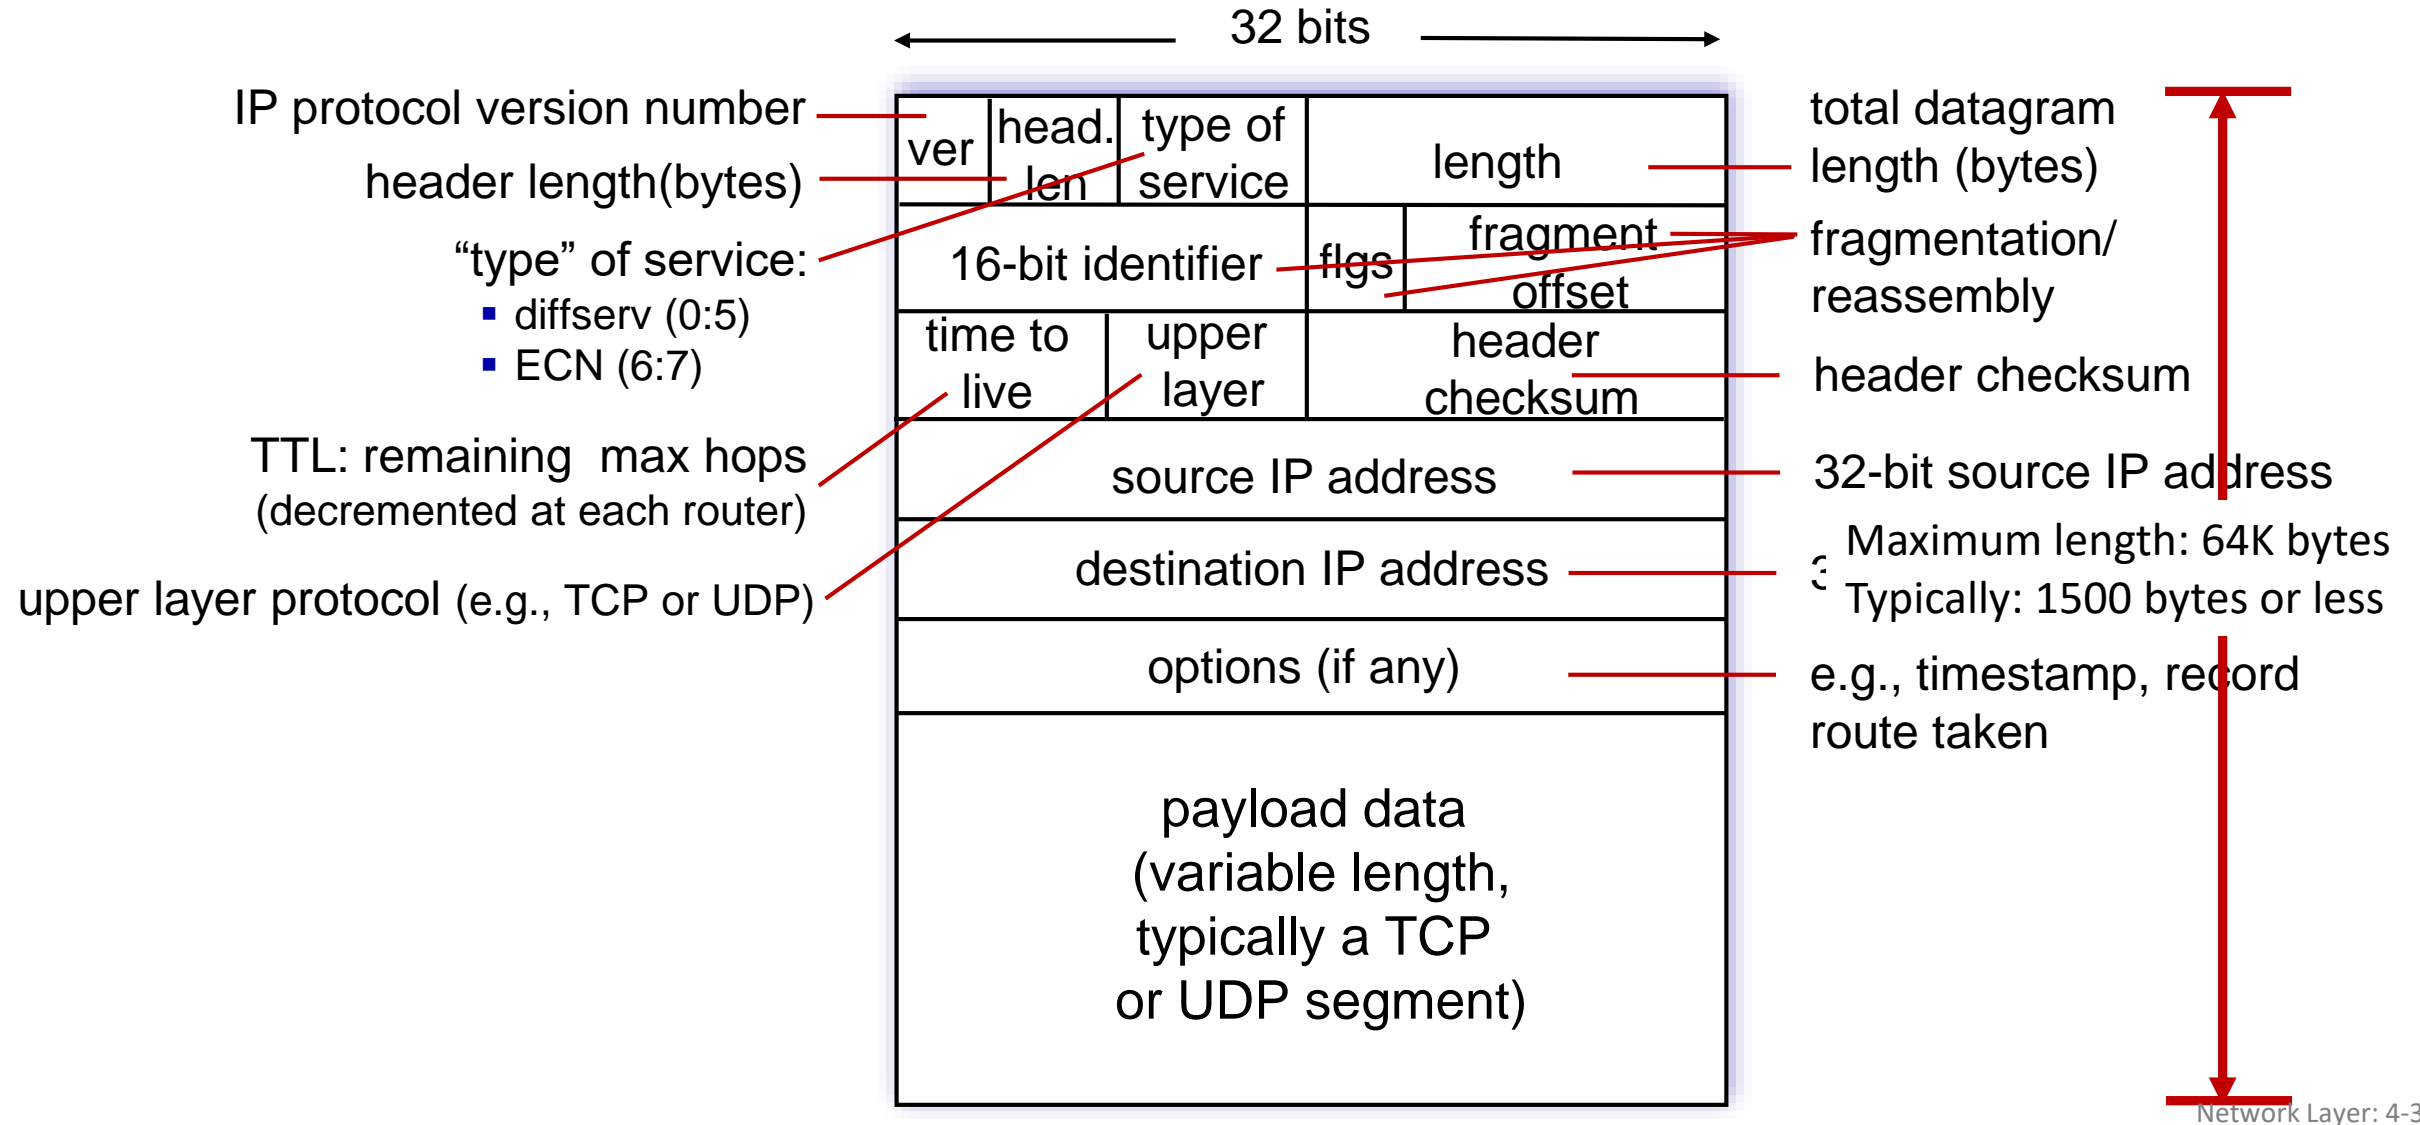
~~4.5 Middleboxes~~

Network Layer: Internet

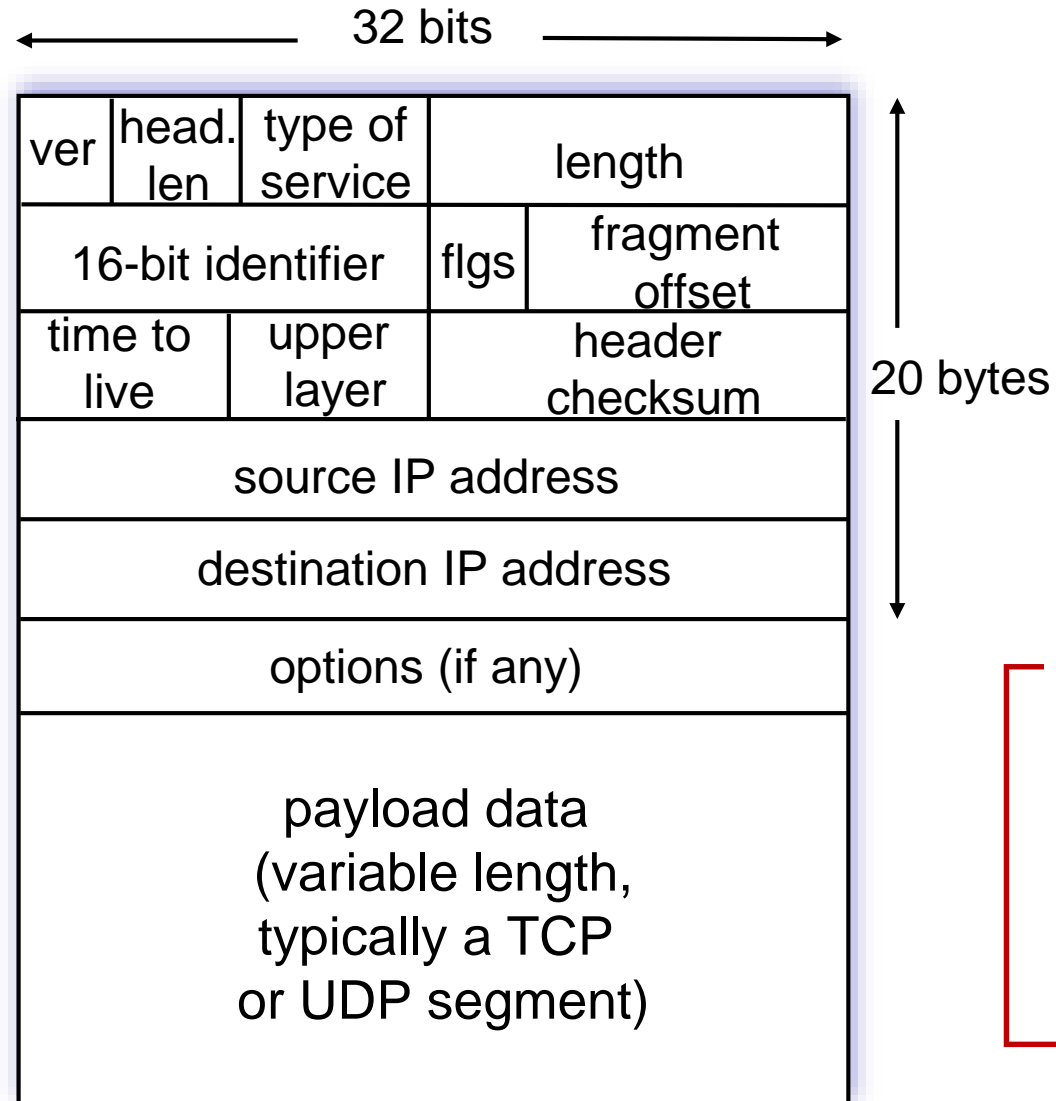
host, router network layer functions:



IP Datagram format



IP Datagram format



- overhead**
- 20 bytes of TCP
 - 20 bytes of IP
 - = 40 bytes + app layer overhead for TCP+IP

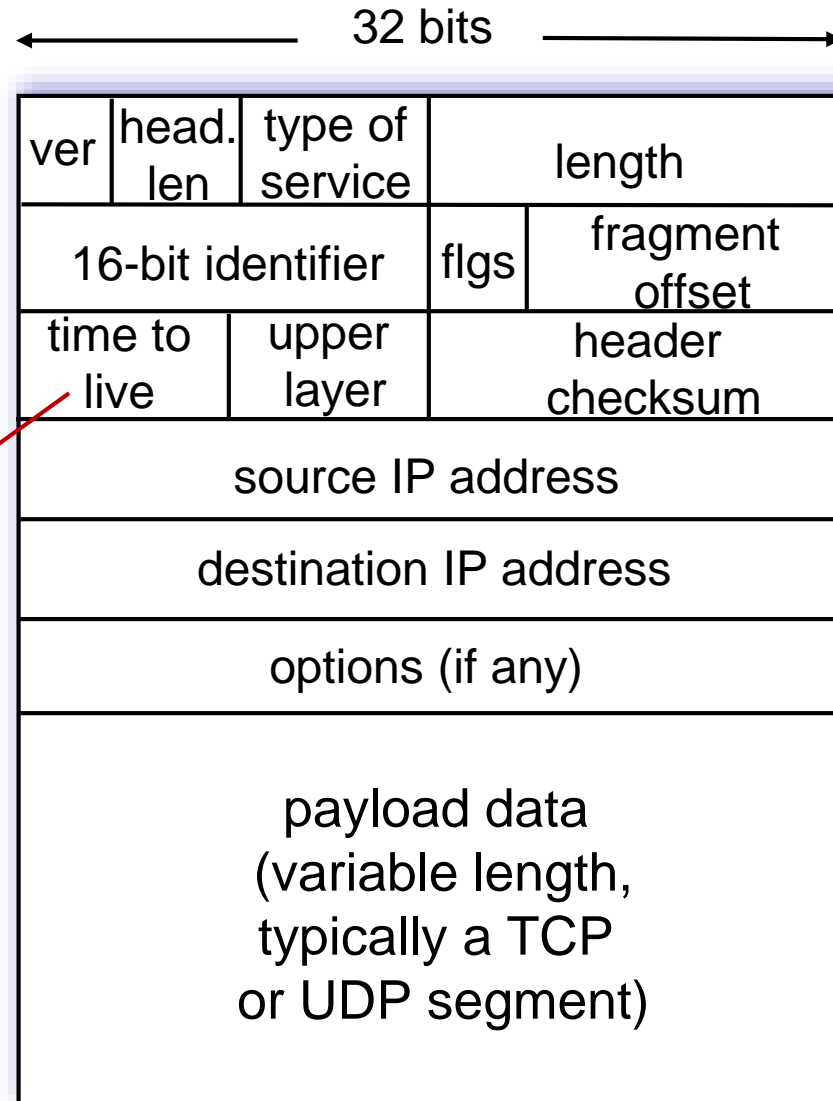
IP Datagram format - TTL

TTL: remaining max hops
(decremented at each router)


TTL

- Purpose:
get rid of looping datagram
- Router process:

```
TTL -= 1
if TTL == 0:
    drop(datagram)
```



TTL and Traceroute



- source sends series of UDP segments to destination
 - 1st set has TTL =1, discarded at 1st router → notify source, source record router and round-trip time.
 - 2nd set has TTL=2, discarded at 2nd router → notify source, source record router and round-trip time.
 - ...
 - reach destination with unlikely port number: destination host notify source with “*port unreachable*”, source record destination and round-trip time.
- Source now has
 - A list of all routers and their round-trip time along the path
 - the destination and destination round-trip time.

```
C:\WINDOWS\system32\cmd.exe
C:\>tracert speedguide.net

Tracing route to speedguide.net [66.230.207.58]
over a maximum of 30 hops:
  0  <1 ms    <1 ms    <1 ms    192.168.1.1
  1  8 ms     *        8 ms     68.152.180.4
  2  8 ms     7 ms     8 ms     68.152.181.197
  3  15 ms    14 ms    14 ms    ixc01jax-ge-1-0-7.bellsouth.net [205.152.70.146]
  4  13 ms    13 ms    13 ms    ber01gnv-pos-1-0-0.bellsouth.net [65.83.239.129]
  5  14 ms    14 ms    13 ms    axr01nsy-so-7-3-0.bellsouth.net [65.83.237.181]
  6  14 ms    13 ms    14 ms    axr00asn-0-3-0.bellsouth.net [65.83.236.121]
  7  13 ms    13 ms    14 ms    65.83.238.142
  8  25 ms    25 ms    25 ms    cr2.ornfl1.ip.att.net [12.122.143.214]
  9  25 ms    25 ms    24 ms    cr1.attga.ip.att.net [12.122.5.142]
 10  24 ms    24 ms    25 ms    ggr7.attga.ip.att.net [12.122.87.61]
 11  25 ms    24 ms    24 ms    192.205.35.214
 12  25 ms    31 ms    36 ms    ae-62-51.ebr2.atlanta2.level3.net [4.68.103.291]
 13  43 ms    42 ms    43 ms    ae-1-6.bar1.tampa1.level3.net [4.69.137.113]
 14  44 ms    43 ms    42 ms    ae-5-5.car1.tampa1.level3.net [4.69.133.49]
 15  43 ms    43 ms    43 ms    ae-14-14.car4.tampa1.level3.net [4.69.133.58]
 16  43 ms    42 ms    43 ms    hostway-cor.car4.tampa1.level3.net [4.71.2.141]
 17  43 ms    43 ms    43 ms    te49.dr5.as30217.net [84.40.24.82]
 18  43 ms    42 ms    43 ms    bce0-ss1.sria.as30217.net [84.40.24.134]
 19  42 ms    43 ms    42 ms    speedguide.net [66.230.207.58]

Trace complete.
```

Network layer: “data plane” roadmap

4.1 Network layer: overview

- data plane
- control plane

4.2 What’s inside a router

- input ports, switching, output ports
- buffer management, scheduling

4.3 IP: the Internet Protocol

- datagram format
- addressing
- network address translation
- IPv6




~~4.4 Generalized Forwarding, SDN~~

- ~~• Match+action~~
- ~~• OpenFlow: match+action in action~~

~~4.5 Middleboxes~~


IP address: introduction

- **IP address:** 32-bit identifier associated with each host or router *interface*
- **interface:** connection between host/router and physical link
 - router's typically have multiple interfaces
 - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)
 - only one interface is active at a time



IP address: decimal → binary

223.1.10.5 = 11011111.01100100.00001010.00000101



223

100

10

5

1. Write down binary table:

7	6	5	4	3	2	1	0		ⁱ
128	64	32	16	8	4	2	1		2 ⁱ

2. convert 10 to binary: 10-8=2; 2-2=0. (deduct numbers: high to low)

3.

0	0	0	0	1	0	1	0
---	---	---	---	---	---	---	---

2. converting 100 to binary: deduct numbers: high to low

100 - 64 = 36; 36 - 32 = 4; 4 - 4 = 0


3.

0	1	1	0	0	1	0	0
---	---	---	---	---	---	---	---

Try 223?

IP address: binary → decimal

223.1.10.5 = 11011111.01100100.00001010.00000101



223

100

10

5

1. Write down binary table:

7	6	5	4	3	2	1	0		ⁱ
128	64	32	16	8	4	2	1		2 ⁱ

2. write binary byte under the table:

0	1	1	0	0	1	0	0
---	---	---	---	---	---	---	---

3. add 1's values:

64 + 32 +

4


= 100

IP addressing: introduction

Q: how are interfaces actually connected?

A: we'll learn about that in chapters 6, 7

A: wired Ethernet interfaces connected by Ethernet switches



A: wireless WiFi interfaces connected by WiFi base station

Subnets

■ *What's a subnet ?*


- device interfaces that can physically reach each other **without passing through an intervening router**
- Analogy: Houses in a **Street**

■ IP addresses have structure:

- **subnet part**: common high order bits
- **host part**: low order bits
- IP address: subnet.**host#**
 223.1.1.**3**
- Street addr: street.**house#**
 Smith St.**3**



3 Smith St




network consisting of 3 subnets

Subnets

Recipe for defining subnets:


- group IP addresses with the same high order part
- creating “islands” of isolated networks
- each isolated network is called a *subnet*
- Write IP address
 - IP address: 223.1.1.3
 - subnet mask: 255.255.255.0
 - shorthand: 223.1.1.3/24



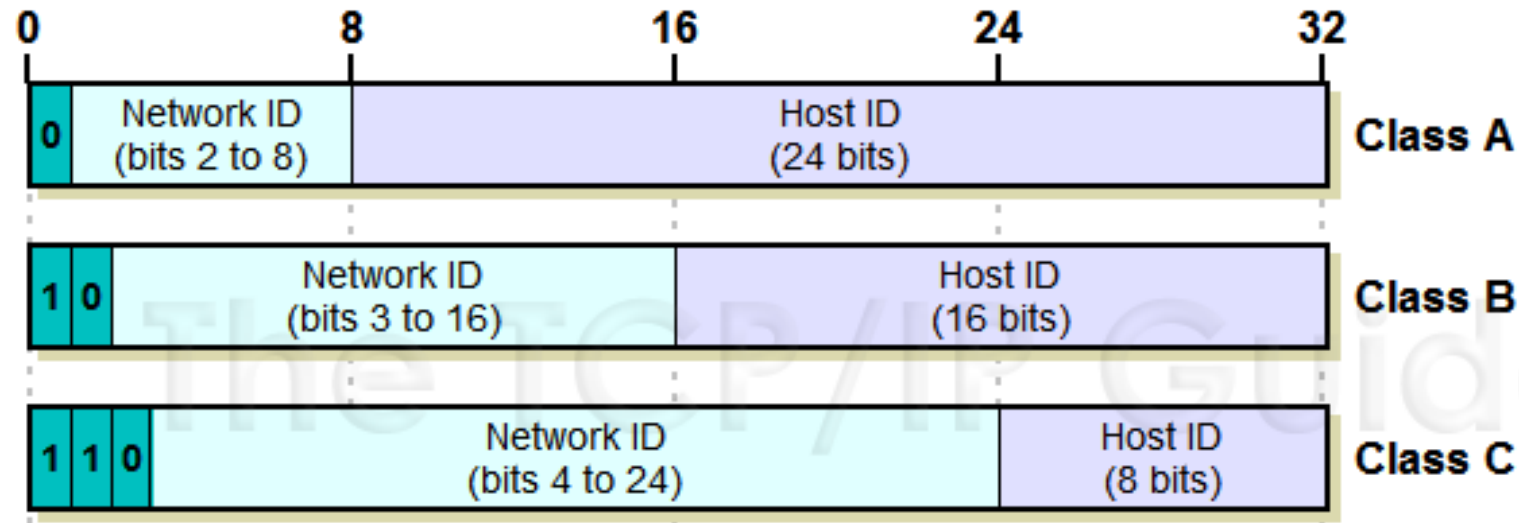
subnet mask: /24
(high-order 24 bits: subnet part of IP address)

Subnets

- where are the subnets?
 - how many subnets?
- what are the subnet addresses?
 - 223.1.1.0/24
 - ...
 - 223.1.9.0/24
- host IP address?
- router IP address?
 - router: gateway to outside



IP Address - Classful




Class	Start Address	End Address	Subnet Mask	shorthand
A	0.0.0.0	127.255.255.255	255.0.0.0	/8
B	128.0.0.0	191.255.255.255	255.255.0.0	/16
C	192.0.0.0	223.255.255.255	255.255.255.0	/24

IP addressing: CIDR

CIDR: Classless **I**nter**D**omain **R**outing (pronounced “cider”)

- subnet portion of address of arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address



Chapter 4 – part1: Summary

- Network layer: overview
- What's inside a router
- IP addressing, subnet
- Forwarding/routing
- Middleboxes



Question: how are forwarding tables (destination-based forwarding) or flow tables (generalized forwarding) computed?

Answer: by the control plane (next chapter)

Lecture done ✓

- Q & A

